

# Investigating objects and data patterns using base R

## EDUC 260A: Managing and Manipulating Data Using R

# Lecture outline

1. Investigate objects, base R
  - 1.1 Functions to describe objects
  - 1.2 Variables names
  - 1.3 View and print data
  - 1.4 Missing values
2. Subsetting using subset operators
  - 2.1 Subset atomic vectors using []
  - 2.2 Subsetting lists/data frames using []
  - 2.3 Subsetting lists/data frames using [[]] and \$
  - 2.4 Subset Data frames by combining [] and \$
3. Subset using subset() function
4. Creating variables
5. Appendix
  - 5.1 Sorting data

## 1 Investigate objects, base R

## Load .Rdata data frames we will use today

Data on off-campus recruiting events by public universities

► Data frame object `df_event`

► One observation per university, recruiting event

► Data frame object `df_school`

► One observation per high school (visited and non-visited)

```
rm(list = ls()) # remove all objects in current environment
```

```
getwd()
```

```
#> [1] "/Users/jaquette/Documents/rclass1/lectures/patterns_base_r"
```

```
#load dataset with one obs per recruiting event
```

```
load(url("https://github.com/ozanj/rclass/raw/master/data/recruiting/recruit_ev
```

```
#load("../..data/recruiting/recruit_event_somevars.Rdata")
```

```
#load dataset with one obs per high school
```

```
load(url("https://github.com/ozanj/rclass/raw/master/data/recruiting/recruit_sc
```

```
#load("../..data/recruiting/recruit_school_somevars.Rdata")
```

## 1.1 Functions to describe objects

## Simple base R functions to describe objects

This section introduces some base R functions to describe objects (some of these you have seen before)

- ▶ list objects, `list.files()` and `ls()`
- ▶ remove objects, `rm()`
- ▶ object type, `typeof()`
- ▶ object length (number of elements), `length()`
- ▶ object structure, `str()`
- ▶ number of rows and columns, `ncol()` and `nrow()`

I use the functions `typeof()`, `length()`, `str()` anytime I encounter a new object

- ▶ Helps me understand the object before I start working with it

# Listing objects

## Files in your working directory

`list.files()` function lists files in your current working directory

► if you run this code from .Rmd file, working directory is location .Rmd file is stored

```
getwd() # what is your current working directory
#> [1] "/Users/jaquette/Documents/rclass1/lectures/patterns_base_r"
list.files()
#> [1] "base_r_week1_video_lecture_script.R" "fp1.JPG"
#> [3] "fp2.JPG" "one_carriage_train_vs_contents.pn
#> [5] "patterns_base_r.pdf" "patterns_base_r.Rmd"
#> [7] "smaller_trains.png" "test.txt"
#> [9] "three_carriage_train.png" "transform-logical.png"
```

# Objects currently open in your R session

## Listing objects currently open in your R session

`ls()` function lists objects currently open in R

```
x <- "hello!"  
ls() # Objects open in R  
#> [1] "df_event" "df_school" "x"
```

## Removing objects currently open in your R session

`rm()` function removes specified objects open in R

```
rm(x)  
ls()  
#> [1] "df_event" "df_school"
```

Command to remove all objects open in R (I don't run it)

```
#rm(list = ls())
```



## Base R functions to describe objects, `typeof()`

`typeof()` function determines the the internal storage type of an object (e.g., logical vector, integer vector, list)

- ▶ syntax

- ▶ `typeof(x)`

- ▶ arguments

- ▶ `x` : any R object

- ▶ help:

`?typeof`

### Examples

- ▶ Recall that a data frame is an object where **type** is a list

```
typeof(c(TRUE,TRUE,FALSE,NA))
```

```
#> [1] "logical"
```

```
typeof(df_event)
```

```
#> [1] "list"
```

```
typeof(x = df_event)
```

```
#> [1] "list"
```

## Base R functions to describe objects, `length()`

`length()` function determines the length of an R object

- ▶ for atomic vectors and lists, `length()` is the number of elements in the object
- ▶ syntax
  - ▶ `length(x)`
- ▶ arguments
  - ▶ `x` : any R object
- ▶ help:

`?length`

Example, length of an atomic vector is

```
length(c(TRUE,TRUE,FALSE,NA))  
#> [1] 4
```

Example, length of a list or data frame

- ▶ length of a list is the number of elements
- ▶ data frame is a list
- ▶ length of a data frame = number of elements = number of variables

```
length(df_event) # = num elements = num columns  
#> [1] 33
```

## Base R functions to describe objects, `str()`

`str()` function compactly displays the structure of an R object

- ▶ “structure” includes type, length, and attribute of object and also nested objects
- ▶ syntax: `str(object)`
- ▶ arguments (partial)
  - ▶ `object` : any R object
  - ▶ `max.level` : max level of nesting to display nested structures; default `NA` = all levels
- ▶ help: `?str`

Example, atomic vectors

```
str(c(TRUE,TRUE,FALSE,NA))  
#> logi [1:4] TRUE TRUE FALSE NA  
str(object = c(TRUE,TRUE,FALSE,NA))  
#> logi [1:4] TRUE TRUE FALSE NA
```

Example, lists/data frames (output omitted)

```
x <- list(c(1,2), list("apple", "orange"), list(2, 3)) # list  
str(x)  
  
str(df_event) # data frame
```

## Base R functions to describe objects, `ncol()` and `nrow()`

`ncol()` `nrow()` and `dim()` functions

### ► Description

► `ncol()` = number of columns; `nrow()` = number of rows

### ► syntax: `ncol(x)` `nrow(x)` `dim(x)`

### ► arguments

► `x`: a vector, array, data frame, or NULL

### ► value/return:

► if object `x` is an atomic vector: `ncol()` and `nrow()` returns NULL

► if object `x` is a list but not a data frame: `ncol()` and `nrow()` returns NULL

► if object `x` is a data frame: `ncol()` and `nrow()` returns integer of length 1

Example, object is a data frame

```
ncol(df_event) # num columns = num elements = num variables
#> [1] 33
nrow(df_event) # num rows = num observations
#> [1] 18680
# can wrap ncol() or nrow() within str() to see what functions return
#str(ncol(df_event))
```

Example, object is atomic vector or list that is not a data frame (output omitted)

```
ncol(c(TRUE,TRUE,FALSE,NA)) # atomic vector
x <- list(c(1,2), list("apple", "orange"), list(2, 3)) # list
nrow(x)
```

## Base R functions to describe objects, `dim()`

`dim()` function returns the dimensions of an object (e.g., number of rows and columns)

- ▶ syntax: `dim(x)`
- ▶ arguments
  - ▶ `x` : a vector, array, data frame, or NULL
- ▶ value/return:
  - ▶ if object `x` is a data frame: `dim()` returns integer of length 2
    - ▶ first element = number of rows; second element = number of columns
  - ▶ if object `x` is an atomic vector: `dim()` returns `NULL`
  - ▶ if object `x` is a list but not a data frame: `dim()` returns `NULL`

Example, object is a data frame

```
dim(df_event) # shows number rows by columns
#> [1] 18680 33

str(dim(df_event)) # can wrap dim() within str() to see what functions return
#> int [1:2] 18680 33
```

Example, object is atomic vector or list that is not a data frame (output omitted)

```
dim(c(TRUE,TRUE,FALSE,NA)) # atomic vector
x <- list(c(1,2), list("apple", "orange"), list(2, 3)) # list
dim(x)
```

## 1.2 Variables names

## names() function

names() function gets or sets the names of elements of an object

### ► syntax:

► get the names of an object: `names(x)`

► set the names of an object: `names(x) <- value`

### ► arguments (partial)

► `x` : an R object

► `value` : a character vector with same length as object `x` or `NULL`

### ► value/return

► `names(x)` returns a character vector of length = `length(x)` in which each element is the name of the element of `x`

Example, get names (of atomic vector)

```
a <- c(v1=1,v2=2,3,v4="hi!") # named atomic vector
```

```
a
```

```
#>      v1      v2      v4
```

```
#>  "1"   "2"   "3"  "hi!"
```

```
length(a)
```

```
#> [1] 4
```

```
names(a)
```

```
#> [1] "v1" "v2" ""  "v4"
```

```
length(names(a)) # investigate length of object names(a)
```

```
#> [1] 4
```

```
str(names(a)) # investigate structure of object names(a)
```

```
#> chr [1:4] "v1" "v2" "" "v4"
```

## names() function

names() function gets or sets the names of elements of an object

### ► syntax:

► get the names of an object: `names(x)`

► set the names of an object: `names(x) <- value`

### ► arguments (partial)

► `x` : an R object

► `value` : a character vector with same length as object `x` or `NULL`

### ► value/return

► `names(x)` returns a character vector of length = `length(x)` in which each element is the name of the element of `x`

Example, set names (of atomic vector)

```
names(a) <- NULL # set names of vector a to NULL
```

```
a
```

```
#> [1] "1" "2" "3" "hi!"
```

```
names(a)
```

```
#> NULL
```

```
names(a) <- c("var1","var2","var3","var4") # set names of vector a
```

```
a
```

```
#> var1 var2 var3 var4
```

```
#> "1" "2" "3" "hi!"
```

```
names(a)
```

```
#> [1] "var1" "var2" "var3" "var4"
```



## Applying `names()` function to a data frame

Recall that a data frame is an object where **type** is a **list** and each **element** is **named**

- ▶ each element is a variable
- ▶ each element name is a variable name

Example (output omitted)

```
names(df_event)
```

Investigate the object `names(df_event)`

```
typeof(names(df_event)) # type = character vector
#> [1] "character"
length(names(df_event)) # length = number of variables in data frame
#> [1] 33
str(names(df_event)) # structure of names(df_event)
#> chr [1:33] "instnm" "univ_id" "instst" "pid" "event_date" "event_type" ...
```

We can even assign a new object based on `names(df_event)`

```
names_event <- names(df_event)
typeof(names_event) # type = character vector
#> [1] "character"
length(names_event) # length = number of variables in data frame
#> [1] 33
str(names_event) # structure of names(df_event)
#> chr [1:33] "instnm" "univ_id" "instst" "pid" "event_date" "event_type" ...
```

## Variable names

Refer to specific named elements of an object using this syntax:

▶ `object_name$element_name`

When object is data frame, refer to specific variables using this syntax:

▶ `data_frame_name$varname`

▶ **This approach to isolating variables is very useful for investigating data**

```
#df_event$instnm  
typeof(df_event$instnm)  
#> [1] "character"  
typeof(df_event$med_inc)  
#> [1] "double"
```

## Variable names

Data frames are lists with the following criteria:

- ▶ each element of the list is (usually) a vector; each element of list is a variable
- ▶ length of data frame = number of variables

```
length(df_event)
```

```
#> [1] 33
```

```
nrow(df_event)
```

```
#> [1] 18680
```

```
#str(df_event)
```

- ▶ each element of the list (i.e., variable) has the same length
  - ▶ Length of each variable is equal to number of observations in data frame

```
typeof(df_event$event_state)
```

```
#> [1] "character"
```

```
length(df_event$event_state)
```

```
#> [1] 18680
```

```
str(df_event$event_state)
```

```
#> chr [1:18680] "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" ..
```

```
typeof(df_event$med_inc)
```

```
#> [1] "double"
```

```
length(df_event$med_inc)
```

```
#> [1] 18680
```

```
str(df_event$med_inc)
```

```
#> num [1:18680] 71714 89122 70136 70136 71024 ...
```

## Variable names

The object `df_school` has one obs per high school

- ▶ variable `visits_by_100751` shows number the of visits by University of Alabama to each high school
- ▶ like all variables in a data frame, the var `visits_by_100751` is just a vector

```
typeof(df_school$visits_by_100751)
#> [1] "integer"
length(df_school$visits_by_100751) # num elements in vector = num obs
#> [1] 21301
str(df_school$visits_by_100751)
#> int [1:21301] 0 0 0 0 0 0 0 0 0 0 0 ...
sum(df_school$visits_by_100751) # sum of values of var across all obs
#> [1] 3338
```

We perform calculations on a variable like we would on any vector of same type

```
v <- c(2,4,6)
typeof(v)
#> [1] "double"
length(v)
#> [1] 3
sum(v)
#> [1] 12
```

## 1.3 View and print data

# Viewing and printing, data frames

Many ways to view/print a data frame object. Here are three ways:

1. Simply type the object name (output omitted)

- ▶ number of observations and rows printed depend on YAML header settings and on object “attributes” (attributes discussed in future week)

```
df_event
```

2. Use the `View()` function to view data in a browser

```
View(df_event)
```

3. `head()` to show the first  $n$  rows. The default is 6 rows.

```
##?head  
#head(df_event)  
head(df_event, n=5)
```

## Viewing and printing, data frames

`obj_name[<rows>,<cols>]` to print specific rows and columns of data frame

- ▶ particularly powerful when combined with sequences (e.g., `1:10` )

Examples (output omitted):

- ▶ Print first five rows, all vars

```
df_event[1:5, ]
```

- ▶ Print first five rows and first three columns

```
df_event[1:5, 1:3]
```

- ▶ Print first three columns of the 100th observation

```
df_event[100, 1:3]
```

- ▶ Print the 50th observation, all variables

```
df_event[50,]
```

## Viewing and printing, variables within data frames

Recall that:

- ▶ `obj_name$var_name` print specific elements (i.e., variables) of a data frame

```
df_event$zip
```

- ▶ each element (i.e., variable) of data frame is an **atomic vector** with **length** = number of observations

```
typeof(df_event$zip)
#> [1] "character"
length(df_event$zip)
#> [1] 18680
```

- ▶ each element of a variable is the value of the variable for one observation

Print specific elements (i.e., observations) of variable based on element position

- ▶ syntax: `obj_name$var_name[<element position>]`
- ▶ vectors don't have "rows" or "columns"; they just have elements
- ▶ syntax combined with sequences (e.g., print first 10 observations)

```
df_event$event_state[1:10] # print obs 1-10 of variable "event_state"
#> [1] "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA"
df_event$event_type[6:10] # print obs 6-10 of variable "event_type"
#> [1] "private hs" "private hs" "public hs" "private hs" "public hs"
```



## Viewing and printing, variables within data frames

Print specific elements (i.e., observations) of variable based on element position

► syntax: `obj_name$var_name[<element position>]`

Example, print individual elements

```
df_event$zip[1:5] # print obs 1-5 of variable for event zip code
#> [1] "01002" "01007" "01020" "01020" "01027"
df_event$zip[1] # print obs 1 of variable for event zip code
#> [1] "01002"
df_event$zip[5] # print obs 5 of variable for event zip code
#> [1] "01027"
df_event$zip[c(1,3,5)] # print obs 5 of variable for event zip code
#> [1] "01002" "01020" "01027"
```

Print specific elements of multiple variables using combine function `c()`

► syntax:

`c(obj_name$var1_name[<element position>], obj_name$var2_name[<element position>])`

► Example: print first five observations of variables "event\_state" and "event\_type"

```
c(df_event$event_state[1:5], df_event$event_type[1:5])
#> [1] "MA" "MA" "MA" "MA" "MA" "public hs"
#> [7] "public hs" "public hs" "public hs" "public hs"
```

## Exercise

Printing exercise using the `df_school` data frame

1. Use the `obj_name[<rows>,<cols>]` syntax to print the first 5 rows and 3 columns of the `df_school` data frame
2. Use the `head()` function to print the first 4 observations
3. Use the `obj_name$var_name[1:10]` syntax to print the first 10 observations of a variable in the `df_school` data frame
4. Use `combine()` to print the first 3 observations of variables "school\_type" & "name"

## Solution

1. Use the `obj_name[<rows>,<cols>]` syntax to print the first 5 rows and 3 columns of the `df_school` data frame

```
df_school[1:5,1:3]
#>   state_code school_type    ncessch
#> 1      AK      public 020000100208
#> 2      AK      public 020000100211
#> 3      AK      public 020000100212
#> 4      AK      public 020000100213
#> 5      AK      public 020000300216
```

## Solution

2. Use the `head()` function to print the first 4 observations

```
head(df_school, n=4)
```

```
#>   state_code school_type      ncessch      name
#> 1         AK      public 020000100208 Bethel Regional High School
#> 2         AK      public 020000100211   Ayagina'ar Elitnaurvik
#> 3         AK      public 020000100212   Kwigillingok School
#> 4         AK      public 020000100213 Nelson Island Area School
#>               address      city zip_code pct_white pct_black
#> 1 1006 Ron Edwards Memorial Dr    Bethel   99559   11.7764    0.5988
#> 2      106 Village Road Kongiganak   99559    0.0000    0.0000
#> 3      108 Village Road Kwigillingok   99622    0.0000    0.0000
#> 4      118 Village Road Toksook Bay   99637    0.0000    0.0000
#>   pct_hispanic pct_asian pct_amerindian pct_other num_fr_lunch total_students
#> 1      1.5968    0.998      84.6307    0.3992      362          501
#> 2      0.0000    0.000      99.4505    0.5495      182          182
#> 3      0.0000    0.000     100.0000    0.0000      116          120
#> 4      0.0000    0.000     100.0000    0.0000      187          201
#>   num_took_math num_prof_math num_took_rla num_prof_rla avgmedian_inc_2564
#> 1          146        24.82          147        24.99        76160.0
#> 2           17         1.70           17         1.70        76160.0
#> 3           14         3.50           14         3.50             NA
#> 4           30         3.00           30         3.00        57656.5
#>   visits_by_110635 visits_by_126614 visits_by_100751 inst_110635 inst_126614
#> 1                0                0                0          CA          CO
#> 2                0                0                0          CA          CO
#> 3                0                0                0          CA          CO
```

## Solution

3. Use the `obj_name$var_name[1:10]` syntax to print the first 10 observations of a variable in the `df_school` data frame

```
df_school$name[1:10]
#> [1] "Bethel Regional High School" "Ayagina'ar Elitnaurvik"
#> [3] "Kwigillingok School"      "Nelson Island Area School"
#> [5] "Alakanuk School"          "Emmonak School"
#> [7] "Hooper Bay School"        "Ignatius Beans School"
#> [9] "Pilot Station School"     "Kotlik School"
```

## Solution

4. Use `combine()` to print the first 3 observations of variables “school\_type” & “name”

```
c(df_school$school_type[1:3],df_school$name[1:3])  
#> [1] "public" "public"  
#> [3] "public" "Bethel Regional High School"  
#> [5] "Ayagina'ar Elitnaurvik" "Kwigillingok School"
```

## 1.4 Missing values

## Missing values

Missing values have the value `NA`

► `NA` is a special keyword, not the same as the character string `"NA"`

use `is.na()` function to determine if a value is missing

► `is.na()` returns a logical vector

```
is.na(5)
#> [1] FALSE
is.na(NA)
#> [1] TRUE
is.na("NA")
#> [1] FALSE
typeof(is.na("NA")) # example of a logical vector
#> [1] "logical"

nvector <- c(10,5,NA)
is.na(nvector)
#> [1] FALSE FALSE TRUE
typeof(is.na(nvector)) # example of a logical vector
#> [1] "logical"

svector <- c("e","f",NA,"NA")
is.na(svector)
#> [1] FALSE FALSE TRUE FALSE
```



# Missing values are “contagious”

What does “contagious” mean?

- ▶ operations involving a missing value will yield a missing value

```
7>5
#> [1] TRUE
7>NA
#> [1] NA
sum(1,2,NA)
#> [1] NA
0==NA
#> [1] NA
2*c(0,1,2,NA)
#> [1] 0 2 4 NA
NA*c(0,1,2,NA)
#> [1] NA NA NA NA
```

## Functions and missing values example, `table()`

`table()` function is useful for investigating categorical variables

```
str(df_event$event_type)
#> chr [1:18680] "public hs" "public hs" "public hs" "public hs" "public hs" ..
table(df_event$event_type)
#>
#> 2yr college 4yr college      other private hs    public hs
#>          951          531        2001         3774        11423
```

## Functions and missing values example, `table()`

By default `table()` ignores `NA` values

```
##?table
str(df_event$school_type_pri)
#>   int [1:18680] NA NA NA NA 1 1 NA 1 NA ...
table(df_event$school_type_pri)
#>
#>      1      2      5
#> 3765      8      1
```

`useNA` argument controls if table includes counts of `NA` s. Allowed values:

- ▶ never ("no") [DEFAULT VALUE]
- ▶ only if count is positive ("ifany");
- ▶ even for zero counts ("always")

```
nrow(df_event)
#> [1] 18680
table(df_event$school_type_pri, useNA="always")
#>
#>      1      2      5 <NA>
#> 3765      8      1 14906
```

Broader point: Most functions that create descriptive statistics have options about how to treat missing values'

- ▶ When investigating data, good practice to *always* show missing values

## 2 Subsetting using subset operators

# Subsetting to Extract Elements

“Subsetting” refers to isolating particular elements of an object

Subsetting operators can be used to select/exclude elements (e.g., variables, observations)

- ▶ there are three subsetting operators: `[]` , `$` , `[]`
- ▶ these operators function differently based on vector types (e.g, atomic vectors, lists, data frames)

## Wichham refers to number of “dimensions” in R objects

An atomic vector is a 1-dimensional object that contains n elements

```
x <- c(1.1, 2.2, 3.3, 4.4, 5.5)
str(x)
#>  num [1:5] 1.1 2.2 3.3 4.4 5.5
```

Lists are multi-dimensional objects

- ▶ Contains n elements; each element may contain a 1-dimensional atomic vector or a multi-dimensional list. Below list contains 3 dimensions

```
list <- list(c(1,2), list("apple", "orange"))
str(list)
#> List of 2
#> $ : num [1:2] 1 2
#> $ :List of 2
#> ..$ : chr "apple"
#> ..$ : chr "orange"
```

Data frames are 2-dimensional lists

- ▶ each element is a variable (dimension=columns)
- ▶ within each variable, each element is an observation (dimension=rows)

```
ncol(df_school)
#> [1] 26
nrow(df_school)
#> [1] 21301
```

## 2.1 Subset atomic vectors using `[]`

# Subsetting elements of atomic vectors

“Subsetting” a vector refers to isolating particular elements of a vector

- ▶ I sometimes refer to this as “accessing elements of a vector”
- ▶ subsetting elements of a vector is similar to “filtering” rows of a data-frame
- ▶ `[]` is the subsetting function for vectors

Six ways to subset an atomic vector using `[]`

1. Using positive integers to return elements at specified positions
2. Using negative integers to exclude elements at specified positions
3. Using logicals to return elements where corresponding logical is `TRUE`
4. Empty `[]` returns original vector (useful for dataframes)
5. Zero vector `[0]`, useful for testing data
6. If vector is “named,” use character vectors to return elements with matching names



## 1. Using positive integers to return elements at specified positions (subset atomic vectors using `[]`)

Create atomic vector `x`

```
(x <- c(1.1, 2.2, 3.3, 4.4, 5.5))  
#> [1] 1.1 2.2 3.3 4.4 5.5  
str(x)  
#> num [1:5] 1.1 2.2 3.3 4.4 5.5
```

`[]` is the subsetting function for vectors

► contents inside `[]` can refer to element number (also called “position”).

► e.g., `[3]` refers to contents of 3rd element (or position 3)

```
x[5] #return 5th element  
#> [1] 5.5
```

```
x[c(3, 1)] #return 3rd and 1st element  
#> [1] 3.3 1.1
```

```
x[c(4,4,4)] #return 4th element, 4th element, and 4th element  
#> [1] 4.4 4.4 4.4
```

```
#Return 3rd through 5th element  
x[3:5]  
#> [1] 3.3 4.4 5.5
```

## 2. Using negative integers to exclude elements at specified positions (subset atomic vectors using `[]`)

Before excluding elements based on position, investigate object

```
x
#> [1] 1.1 2.2 3.3 4.4 5.5

length(x)
#> [1] 5

str(x)
#> num [1:5] 1.1 2.2 3.3 4.4 5.5
```

Use negative integers to exclude elements based on element position

```
x[-1] # exclude 1st element
#> [1] 2.2 3.3 4.4 5.5

x[c(3,1)] # 3rd and 1st element
#> [1] 3.3 1.1

x[-c(3,1)] # exclude 3rd and 1st element
#> [1] 2.2 4.4 5.5
```

### 3. Using logicals to return elements where corresponding logical is TRUE (subset atomic vectors using `[]`)

```
x  
#> [1] 1.1 2.2 3.3 4.4 5.5
```

When using `x[y]` to subset `x`, good practice to have `length(x)==length(y)`

```
length(x) # length of vector x  
#> [1] 5  
length(c(TRUE,FALSE,TRUE,FALSE,TRUE)) # length of y  
#> [1] 5  
length(x) == length(c(TRUE,FALSE,TRUE,FALSE,TRUE)) # condition true  
#> [1] TRUE  
x[c(TRUE,TRUE,FALSE,FALSE,TRUE)]  
#> [1] 1.1 2.2 5.5
```

Recycling rules:

- ▶ in `x[y]`, if `x` is different length than `y`, R “recycles” length of shorter to match length of longer

```
length(c(TRUE,FALSE))  
#> [1] 2  
x  
#> [1] 1.1 2.2 3.3 4.4 5.5  
x[c(TRUE,FALSE)]  
#> [1] 1.1 3.3 5.5
```

### 3. Using logicals to return elements where corresponding logical is TRUE (subset atomic vectors using `[]`)

```
x  
#> [1] 1.1 2.2 3.3 4.4 5.5
```

Note that a missing value ( `NA` ) in the index always yields a missing value in the output:

```
x[c(TRUE, FALSE, NA, TRUE, NA)]  
#> [1] 1.1 NA 4.4 NA
```

Return all elements of object `x` where element is greater than 3:

```
x # print object X  
#> [1] 1.1 2.2 3.3 4.4 5.5  
x>3 # for each element of X, print T/F whether element value > 3  
#> [1] FALSE FALSE TRUE TRUE TRUE  
str(x>3)  
#> logi [1:5] FALSE FALSE TRUE TRUE TRUE  
x[x>3] # prints only the values that had TRUE at that position  
#> [1] 3.3 4.4 5.5
```



4. Empty `[]` returns original vector (subset atomic vectors using `[]`)

```
x
#> [1] 1.1 2.2 3.3 4.4 5.5

x[]
#> [1] 1.1 2.2 3.3 4.4 5.5
```

This is useful for sub-setting data frames, as we will show below

## 5. Zero vector [0] (subset atomic vectors using [])

Zero vector, `x[0]`

- ▶ R interprets this as returning element 0

```
x[0]
```

```
#> numeric(0)
```

Wickham states:

- ▶ “This is not something you usually do on purpose, but it can be helpful for generating test data.”

## 6. If vector is named, character vectors to return elements with matching names (subset atomic vectors using `[]`)

Create vector `y` that has values of vector `x` but each element is named

```
x
#> [1] 1.1 2.2 3.3 4.4 5.5

(y <- c(a=1.1, b=2.2, c=3.3, d=4.4, e=5.5))
#>   a    b    c    d    e
#> 1.1 2.2 3.3 4.4 5.5
```

Return elements of vector based on name of element

► enclose element names in single `' '` or double `" "` quotes

```
#show element named "a"
y["a"]
#>   a
#> 1.1

#show elements "a", "b", and "d"
y[c("a", "b", "d" )]
#>   a    b    d
#> 1.1 2.2 4.4
```



## 2.2 Subsetting lists/data frames using []

## Subsetting lists using []

Using `[]` operator to subset lists works the same as subsetting atomic vector

► Using `[]` with a list always returns a list

```
list_a <- list(list(1,2),3,"apple")
str(list_a)
#> List of 3
#> $ :List of 2
#> ..$ : num 1
#> ..$ : num 2
#> $ : num 3
#> $ : chr "apple"

#create new list that consists of elements 3 and 1 of list_a
list_b <- list_a[c(3, 1)]
str(list_b)
#> List of 2
#> $ : chr "apple"
#> $ :List of 2
#> ..$ : num 1
#> ..$ : num 2

#show elements 3 and 1 of object list_a
#str(list_a[c(3, 1)])
```

# Subsetting data frames using `[]`

Recall that a data frame is just a particular kind of list

- ▶ each element = a column = a variable

Using `[]` with a list always returns a list

- ▶ Using `[]` with a data frame always returns a data frame

Two ways to use `[]` to extract elements of a data frame

1. use “single index” `df_name[<columns>]` to extract columns (variables) based on element position number (i.e., column number)
2. use “double index” `df_name[<rows>, <columns>]` to extract particular rows and columns of a data frame

## Subsetting data frames using [] to extract columns (variables) based on element position

Use “single index” `df_name[<columns>]` to extract columns (variables) based on element number (i.e., column number)

Examples [output omitted]

```
names(df_event)
```

```
#extract elements 1 through 4 (elements=columns=variables)
```

```
df_event[1:4]
```

```
df_event[c(1,2,3,4)]
```

```
str(df_event[1:4])
```

```
#extract columns 13 and 7
```

```
df_event[c(13,7)]
```

## Subsetting Data Frames to extract columns (variables) and rows (observations) based on positionality

use “double index” syntax `df_name[<rows>, <columns>]` to extract particular rows and columns of a data frame

► often combined with sequences (e.g., `1:10` )

```
#Return rows 1-3 and columns 1-4
```

```
df_event[1:3, 1:4]
```

```
#>      instnm univ_id instst  pid  
#> 1 UM Amherst  166629    MA 57570  
#> 2 UM Amherst  166629    MA 56984  
#> 3 UM Amherst  166629    MA 57105
```

```
#Return rows 50-52 and columns 10 and 20
```

```
df_event[50:52, c(10,20)]
```

```
#>      event_state pct_tworaces_zip  
#> 50             MA             1.9831  
#> 51             MA             1.9831  
#> 52             MA             1.9831
```

## Subsetting Data Frames to extract columns (variables) and rows (observations) based on positionality

use “double index” syntax `df_name[<rows>, <columns>]` to extract particular rows and columns of a data frame

recall that empty `[]` returns original object (output omitted)

```
#return original data frame
```

```
df_event[]
```

```
#return specific rows and all columns (variables)
```

```
df_event[1:5, ]
```

```
#return all rows and specific columns (variables)
```

```
df_event[, c(1,2,3)]
```

## Use `[]` to extract data frame columns based on variable names

Selecting columns from a data frame by subsetting with `[]` and list of element names (i.e., variable names) enclosed in quotes

“single index” approach extracts specific variables, all rows (output omitted)

```
df_event[c("instnm", "univ_id", "event_state")]
```

“Double index” approach extracts specific variables and specific rows

► syntax `df_name[<rows>, <columns>]`

```
df_event[1:5, c("instnm", "event_state", "event_type")]
```

```
#>      instnm event_state event_type
#> 1  UM Amherst      MA public hs
#> 2  UM Amherst      MA public hs
#> 3  UM Amherst      MA public hs
#> 4  UM Amherst      MA public hs
#> 5 Stony Brook      MA public hs
```

## Student exercises

Use subsetting operators from base R in extracting columns (variables), observations:

1. Use both “single index” and “double index” in subsetting to create a new dataframe by extracting the columns `instnm`, `event_date`, `event_type` from the `df_event` data frame. And show what columns (variables) are in the newly created dataframe.
2. Use subsetting to return rows 1-5 of columns `state_code`, `name`, `address` from the `df_school` data frame.



# Solution to Student Exercises

## Solution to 1

**base R** using subsetting operators

```
# single index
df_event_br <- df_event[c("instnm", "event_date", "event_type")]
#double index
df_event_br <- df_event[, c("instnm", "event_date", "event_type")]
names(df_event_br)
#> [1] "instnm"      "event_date" "event_type"
```

## Solution to 2

**base R** using subsetting operators

```
df_school[1:5, c("state_code", "name", "address")]
#>   state_code      name      address
#> 1      AK Bethel Regional High School 1006 Ron Edwards Memorial Dr
#> 2      AK      Ayagina'ar Elitnaurvik      106 Village Road
#> 3      AK      Kwigillingok School      108 Village Road
#> 4      AK Nelson Island Area School      118 Village Road
#> 5      AK      Alakanuk School      9 School Road
```

## 2.3 Subsetting lists/data frames using `[[` and `$`

## Subset single element from object using `[[ ]]` operator, atomic vectors

So far we have used `[]` to extract elements from an object

- ▶ Apply `[]` to atomic vector: returns atomic vector with elements you requested
- ▶ Apply `[]` to list: returns list with elements you requested

`[[ ]]` also extract elements from an object

- ▶ Applying `[[ ]]` to atomic vector gives same result as `[]`; that is, an atomic vector with element you request

```
(x <- c(1.1, 2.2, 3.3, 4.4, 5.5))
```

```
#> [1] 1.1 2.2 3.3 4.4 5.5
```

```
str(x[3])
```

```
#>  num 3.3
```

```
str(x[[3]])
```

```
#>  num 3.3
```

- ▶ Caveat: when applying `[[ ]]` to atomic vector, you can only subset a single element

```
x[c(3,4)] # single bracket; this works
```

```
#> [1] 3.3 4.4
```

```
#x[[c(3,4)]] # double bracket; this won't work
```

## Subsetting lists using `[]` vs. `[[ ]]`, introduce “train metaphor”

Applying `[[ ]]` to a list

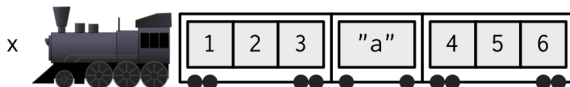
- Understanding what `[]` vs. `[[ ]]` does to a list is very important but requires some explanation!

*Advanced R* [chapter 4.3](#) by Wickham uses the “train metaphor” to explain a list vs. **contents** of a list and how this relates to `[]` vs. `[[ ]]`

Below code chunk makes a list named `list_x` that contains 3 elements

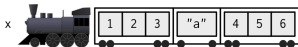
```
list_x <- list(1:3, "a", 4:6) # create list object list_x
```

In our train metaphor, object `list_x` is a train that contains 3 carriages

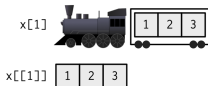


## Subsetting lists using `[]` vs. `[[[]]`, introduce “train metaphor”

list object `list_x` is a train that contains 3 carriages



When we “subset a list” – that is, extract one or more elements from the list – we have two broad choices (image below)



1. Extracting elements using `[]` always returns a list, usually one with fewer elements

► you can think of this as a train with fewer carriages

```
#str(list_x)
#str(list_x[1]) # returns a list
#> List of 1
#> $ : int [1:3] 1 2 3
```

2. Extracting element using `[[[]]` returns **contents** of particular carriage

► I say applying `[[[]]` to a list or data frame returns a simpler object that moves up one level of hierarchy

```
str(list_x[[1]]) # returns an atomic vector
#> int [1:3] 1 2 3
```

## Subset lists using `[]` vs. `[[[]]`, deepen understanding of `[]`

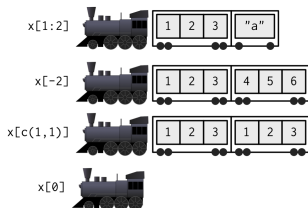
Rules about applying subset operator `[]` to a list

- ▶ Applying `[]` to a list always returns a list
- ▶ Resulting list contains 1 or more elements depending on what typed inside `[]`

Here is a list object named `list_x`

```
list_x <- list(1:3, "a", 4:6)
```

Here is an image of a few “trains” that can be created by applying `[]` to `list_x`



And here is code to create the “trains” shown in above image (output omitted)

```
list_x[1:2]
list_x[-2]
list_x[c(1,1)]
list_x[0]
list_x[] # returns the original list; not shown in above train picture
```

## Subset lists using `[]` vs. `[[[]]`, deepen understanding of `[[[]]`

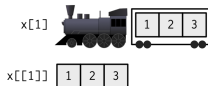
Rules about applying subset operator `[[[]]` to a list

- ▶ Can apply `[[[]]` to return the **contents** of a **single element** of a list

Create list `list_x` and show “train” Image of applying `list_x[1]`

vs. `list_x[[1]]`

```
list_x <- list(1:3, "a", 4:6)
```



Object created by `list_x[1]` is a list with one element (output omitted)

```
list_x[1]  
str(list_x[1])
```

Object created by `list_x[[1]]` is a vector with 3 elements (output omitted)

- ▶ `list_x[[1]]` gives us “contents” of element 1
- ▶ Since element 1 contains a numeric vector, object created by `list_x[[1]]` is a numeric vector

```
list_x[[1]]  
str(list_x[[1]])
```

## Subset lists using `[]` vs. `[[[]]`, deepen understanding of `[[[]]`

Rules about applying subset operator `[[[]]` to a list

- ▶ Can apply `[[[]]` to return the **contents** of a **single element** of a list

```
list_x <- list(1:3, "a", 4:6) # create list list_x
```

We cannot use `[[[]]` to subset multiple elements of a list (output omitted)

- ▶ e.g., we could write `list_x[[2]]` but not `list_x[[2:3]]`

```
list_x[[c(2)]] # this works, subset element 2 using [[[]]  
list_x[[c(2,3)]] # this doesn't work; subset element 2 and 3 using [[[]]  
list_x[c(2,3)] # this works; subset element 2 and 3 using []
```



## Subset lists using `[]` vs. `[[[]]`, deepen understanding of `[[[]]`

Like `[]`, can use `[[[]]` to return contents of **named** elements specified using quotes

► syntax: `obj_name[["element_name"]]`

Same list as before, but this time elements named

```
list_x <- list(var1=1:3, var2="a", var3=4:6)
```

Subset list `list_x` using `[[[]]` element names vs. element position

```
list_x[["var1"]]
#> [1] 1 2 3
# note: list_x[[1]] # same as above
list_x[["var3"]]
#> [1] 4 5 6
```

Can do same thing with data frames because data frames are lists (output omitted)

- e.g., `df_event[["zip"]]` returns contents of element named "zip"
- object created by `df_event[["zip"]]` is character vector of length = 18,680

```
# df_event[["zip"]] # this works but long output
str(df_event[["zip"]]) # character vector of length 18,860
typeof(df_event[["zip"]])
length(df_event[["zip"]])
str(df_event["zip"]) # by contrast, this is a dataframe w/ one variable
```

## General rules of applying `[]` vs `[[] ]` to (nested) objects

What we just learned about applying `[]` vs  `[[] ]` to lists applies more generally to “nested objects”

- ▶ “nested objects” are objects with a hierarchical structure such that an element of an object contains another object

General rules of applying `[]` vs.  `[[] ]` to nested objects

- ▶ subset any object `x` using `[]` will return object with same data structure as `x`
- ▶ subset any object `x` using  `[[] ]` will return an object that may or may not have same data structure of `x`
  - ▶ if object `x` is not a nested object, then applying  `[[] ]` to a single element of `x` will return object with same data structure as `x`
  - ▶ if object `x` has a nested data structure, then then applying  `[[] ]` to a single element of `x` will “move up one level of hierarchy” to extract the **contents** of an element within the object `x`

## Subset lists/data frames using \$

```
list_x <- list(var1=1:3, var2="a", var3=4:6)
```

`obj_name$element_name` is shorthand operator for `obj_name[["element_name"]]`

These three lines of code all give the same result

```
list_x[[1]]  
#> [1] 1 2 3  
list_x[["var1"]]  
#> [1] 1 2 3  
list_x$var1  
#> [1] 1 2 3
```

`df_name$var_name` : easiest way in base R to refer to variable in a data frame

► these two lines of code are equivalent

```
str(df_event[["zip"]])  
#> chr [1:18680] "01002" "01007" "01020" "01020" "01027" "01027" "01027" ...  
str(df_event$zip)  
#> chr [1:18680] "01002" "01007" "01020" "01020" "01027" "01027" "01027" ...
```

## 2.4 Subset Data frames by combining [] and \$

# Subset Data Frames by combining `[]` and `$`, Motivation

## Motivation

- ▶ When working with data frames we often want to isolate those observations that satisfy certain conditions
- ▶ This is often referred to as “filtering”
  - ▶ We filter observations based on the values of one or more variables
- ▶ Perhaps you have seen “filtering” in Microsoft Excel
  - ▶ open some spreadsheet that contains variables (columns) and observations (rows)
  - ▶ click on `Data` » `Filter` and then filter observations based on values of variable(s)

## Filtering example using data frame `df_school`

- ▶ Observations:
  - ▶ One observation per high school (public and private)
- ▶ Variables:
  - ▶ high school characteristics; number of off-campus recruiting visits from particular universities
  - ▶ NCES ID for UC Berkeley is `110635`
  - ▶ variable `visits_by_110635` shows number of visits a high school received from UC Berkeley
- ▶ **Task:**
  - ▶ Isolate observations where the high school received at least 1 visit from UC Berkeley

## Subset Data Frames by combining `[]` and `$`

### Task:

- ▶ Using `df_school`, isolate obs where school received at least 1 visit from UC Berkeley

General syntax: `df_name[df_name$var_name <condition>, ]`

- ▶ Note that syntax uses “double index” `df_name[<rows>, <columns>]` syntax
  - ▶ Therefore, the `<condition>` in above syntax is isolating `<rows>`
- ▶ **Cannot** use “single index” syntax `df_name[<columns>]`

Solution to task (output omitted)

- ▶ Note: below code filters observations but keeps all variables

```
df_school[df_school$visits_by_110635 >= 1, ]
```

## Subset Data Frames by combining `[]` and `$`, decompose syntax

**Task:** Isolate obs where school received at least 1 visit from UC Berkeley

▶ general syntax: `df_name[df_name$var_name <condition>, ]`

▶ solution: `df_school[df_school$visits_by_110635 >= 1, ]`

Decomposing syntax `df_school[df_school$visits_by_110635 >= 1, ]`

▶ `df_school$visits_by_110635 >= 1`: returns a logical ( `TRUE` / `FALSE` ) atomic vector with length equal to number of obs in `df_school`

```
typeof(df_school$visits_by_110635 >= 1)
```

```
length(df_school$visits_by_110635 >= 1)
```

```
str(df_school$visits_by_110635 >= 1)
```

▶ `df_school[df_school$visits_by_110635 >= 1, ]`

▶ uses “double index” `df_name[<rows>, <columns>]` syntax to extract rows, columns

▶ rows: extract rows where `df_school$visits_by_110635 >= 1` is `TRUE`

▶ columns: since `<columns>` is empty, extracts all columns

▶ **key point:** `df_name[df_name$var_name <condition>, ]` is “subset a vector approach #3”: “Using logicals to return elements where condition `TRUE`”

▶ example using atomic vectors (output omitted)

```
x <- c(1.1, 2.2, 3.3, 4.4, 5.5)
```

```
x[x>3]
```

## Subset Data Frames by combining `[]` and `$`, keep desired columns

- ▶ General syntax to filter desired observations (rows) and variables (columns) of data frame:

- ▶ `df_name[df_name$var_name <condition>, <desired columns>]`

### Tasks (output omitted)

- ▶ Extract observations where the high school received at least 1 visit from UC Berkeley and the first three columns

```
df_school[df_school$visits_by_110635 >= 1, 1:3]
```

- ▶ Extract observations where the high school received at least 1 visit from UC Berkeley and variables "state\_code" "school\_type" "name"

```
df_school[df_school$visits_by_110635 >= 1, c("state_code", "school_type", "name")]
```



## Subset Data Frames by combining `[]` and `$`, more examples

Syntax:

- ▶ filter based on one variable:

- ▶ `df_name[df_name$var_name <condition>, <columns>]`

- ▶ Example syntax to filter based on two conditions being true

- ▶ `df_name[df_name$var_name <condition> & df_name$var_name <condition>, <columns>]`

Pro tip:

- ▶ wrap above syntax within `nrow()` function to count how many observations (rows) satisfy the condition (as opposed to printing all rows that satisfy condition)

### Tasks

- ▶ Count obs where high schools received at least 1 visit by Bama (100751) **and** at least one visit by Berkeley (110635)

```
nrow(df_school[df_school$visits_by_110635 >= 1 &
               df_school$visits_by_100751 >= 1, ])
```

*#> [1] 247*  
*# Equivalently:*  
*# nrow(df\_school[df\_school[["visits\_by\_110635"]] >= 1 &*  
*# df\_school[["visits\_by\_100751"]] >= 1, ])*

- ▶ Count obs where schools received 1+ visit by Bama **or** 1+ visit by Berkeley

```
nrow(df_school[df_school$visits_by_110635 >= 1
               | df_school$visits_by_100751 >= 1, ])
```

## Logical operators for comparisons

- Logical operators to isolate/filter observations of data frame

Symbol	Meaning
<code>==</code>	Equal to
<code>!=</code>	Not equal to
<code>&gt;</code>	greater than
<code>&gt;=</code>	greater than or equal to
<code>&lt;</code>	less than
<code>&lt;=</code>	less than or equal to
<code>&amp;</code>	AND
<code> </code>	OR
<code>%in%</code>	includes

- Visualization of “Boolean” operators (e.g., AND, OR, AND NOT)

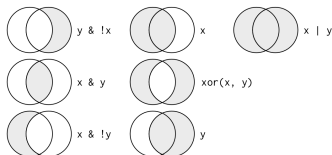


Figure 1: “Boolean” operations, x=left circle, y=right circle, from Wichkam (2018)

## Subset Data Frames by combining `[]` and `$`, more examples

**Example:** Count the number of out-of-state high schools that UC Berkeley visited.

```
# The `inst_110635` variable contains the home state of UC Berkeley
unique(df_school$inst_110635)
#> [1] "CA"

# Filter for schools visited by UC Berkeley AND whose state is not "CA"
nrow(df_school[df_school$visits_by_110635 >= 1 &
               df_school$state_code != df_school$inst_110635, ])
#> [1] 302
```

**Example:** Count the number of schools in the Northeast that received a visit from either UC Berkeley, U of Alabama, or CU Boulder.

```
# Vector containing states located in the Northeast region
northeast_states <- c('CT', 'ME', 'MA', 'NH', 'RI', 'VT', 'NJ', 'NY', 'PA')

# Filter for schools in the Northeast AND visited by any of the 3 univs
nrow(df_school[df_school$state_code %in% northeast_states &
               (df_school$visits_by_110635 >= 1 |
                df_school$visits_by_100751 >= 1 |
                df_school$visits_by_126614 >= 1), ])
#> [1] 544
```

## Subset Data Frames by combining `[]` and `$`, `NA` Observations

Filtering observations of data frame using `[]` combined with `$` is more complicated in the presence of missing values ( `NA` values)

The next few slides will explain

- ▶ why it is more complicated
- ▶ how to filter correctly when `NA` s are present

## Subset Data Frames by combining `[]` and `$`, `NA` Observations

When sub-setting via `[]` combined with `$`, result will include:

- ▶ rows where condition is `TRUE`
- ▶ as well as rows with `NA` (missing) values for condition.

**Task** (using `df_event`, which has one obs per university, recruiting event)

- ▶ How many events at public HS with at least \$50k median household income?

```
sum(is.na(df_event$med_inc)) # number of observations (all events) with missing  
#> [1] 470
```

```
#num obs event_type=="public hs" and med_inc is missing  
nrow(df_event[df_event$event_type == "public hs"  
  & is.na(df_event$med_inc)==1 , ]) # note TRUE evaluates to 1  
#> [1] 75
```

```
#num obs event_type=="public hs" & med_inc is not NA & med_inc >= $50,000  
nrow(df_event[df_event$event_type == "public hs"  
  & is.na(df_event$med_inc)==0 & df_event$med_inc>=50000 , ]) # note FALSE eva  
#> [1] 9941
```

```
#num obs event_type=="public hs" and med_inc >= $50,000  
nrow(df_event[df_event$event_type == "public hs"  
  & df_event$med_inc>=50000 , ])  
#> [1] 10016
```

## Subset Data Frames by combining `[]` and `$`, NA Observations

To exclude rows where condition is NA if subset using `[]` combined w/ `$`

- ▶ use `which()` to ask only for values where condition evaluates to `TRUE`
- ▶ `which()` returns position numbers for elements where condition is `TRUE`

```
#?which  
c(TRUE,FALSE,NA,TRUE)  
#> [1] TRUE FALSE NA TRUE  
str(c(TRUE,FALSE,NA,TRUE))  
#> logi [1:4] TRUE FALSE NA TRUE  
which(c(TRUE,FALSE,NA,TRUE))  
#> [1] 1 4
```

Task: Count events at public HS with at least \$50k median household income?

```
#Base R, `[]` combined with `$`; without which()  
nrow(df_event[df_event$event_type == "public hs" & df_event$med_inc>=50000, ])  
#> [1] 10016  
  
#Base R, `[]` combined with `$`; with which()  
nrow(df_event[which(df_event$event_type == "public hs"  
  & df_event$med_inc>=50000), ])  
#> [1] 9941
```

# Student Exercises

Subsetting Data Frames with `[]` and `$` :

1. Show how many public high schools in California with at least 50% Latinx (hispanic in data) student enrollment from `df_school`.
2. Show how many out-state events at public high schools with more than \$30K median from `df_event` (do not forget to exclude missing values).

## Solution to Student Exercises

Solution to 1

base R using `[]` and `$`

```
df_school_br1<- df_school[df_school$school_type == "public"  
                           & df_school$pct_hispanic >= 50  
                           & df_school$state_code == "CA", ]  
nrow(df_school_br1)  
#> [1] 713
```



## Solution to Student Exercises

Solution to 2:

**base R** using `[]` and `$`

```
# use is.na to exclude NA
nrow(df_event[df_event$event_type == "public hs" & df_event$event_inst == "Out-S
          & df_event$med_inc > 30000 & is.na(df_event$med_inc) == 0, ])
#> [1] 7784

# use which to exclude NA
nrow(df_event[which(df_event$event_type == "public hs" & df_event$event_inst ==
          & df_event$med_inc > 30000 ), ])
#> [1] 7784
```

### 3 Subset using subset() function

## Subset function

The `subset()` is a base R function to “filter” observations from some object `x`

- ▶ object `x` can be a matrix, data frame, list
- ▶ `subset()` automatically excludes elements/rows with `NA` for condition
- ▶ Can also use `subset()` to select variables
- ▶ what `subset()` function returns:
  - ▶ “An object similar to `x` contain just the selected ...rows and columns (for a matrix or data frame)”
- ▶ `subset()` can be combined with:
  - ▶ assignment ( `<-` ) to create new objects
  - ▶ `nrow()` to count number of observations that satisfy criteria

?subset

Syntax [when object is data frame]: **`subset(x, subset, select, drop = FALSE)`**

- ▶ `x` is object to be subset
- ▶ `subset` is the logical expression(s) (evaluates to `TRUE/FALSE` ) indicating elements (rows) to keep
- ▶ `select` indicates columns to select from data frame (if argument is not used default will keep all columns)
- ▶ `drop` to preserve original **dimensions** [SKIP]

## Subset function, examples

Recall the previous example where we count events at public HS with at least \$50k median household income.

► *Note.* `subset()` automatically excludes rows where condition is `NA` :

```
#Base R, `[ ]` combined with `$`, without which(); includes `NA`  
nrow(df_event[df_event$event_type == "public hs"  
              & df_event$med_inc >= 50000, ])  
#> [1] 10016
```

```
#Base R, `[ ]` combined with `$`, with which(); excludes `NA`  
nrow(df_event[which(df_event$event_type == "public hs"  
                   & df_event$med_inc >= 50000), ])  
#> [1] 9941
```

```
#Base R, `subset()`; excludes `NA`  
nrow(subset(df_event, event_type == "public hs"  
            & med_inc >= 50000))  
#> [1] 9941
```

```
#Base R, `subset()`; excludes `NA`; explicitly name arguments of subset()  
nrow(subset(x = df_event, subset = event_type == "public hs"  
            & med_inc >= 50000))  
#> [1] 9941
```

## Subset function, examples

Using `df_school`, show all public high schools that are at least 50% Latinx (var=`pct_hispanic`) student enrollment in California

► Using base R, `subset()` [output omitted]

```
#public high schools with at least 50% Latinx student enrollment
subset(x= df_school, subset = school_type == "public" & pct_hispanic >= 50
      & state_code == "CA")
```

► Can wrap `subset()` within `nrow()` to count number of observations that satisfy criteria

```
nrow(subset(df_school, school_type == "public" & pct_hispanic >= 50
           & state_code == "CA"))
#> [1] 713
```

## Subset function, examples

Note that `subset()` identify the number of observations for which the condition is `TRUE`

```
nrow(subset(x = df_school, subset = TRUE))  
#> [1] 21301  
nrow(subset(x = df_school, subset = FALSE))  
#> [1] 0
```

## Subset function, examples

Count all CA public high schools that are at least 50% Latinx and received at least 1 visit from UC Berkeley (var= visits\_by\_110635 )

```
nrow(subset(df_school, school_type == "public" & pct_hispanic >= 50
  & state_code == "CA" & visits_by_110635 >= 1))
#> [1] 100
```

## Subset function, examples

`subset()` can also use `%in%` operator, which is more efficient version of **OR** operator `|`

- ▶ Count number of schools from MA, ME, or VT that received at least one visit from University of Alabama (var= `visits_by_100751` )

```
nrow(subset(df_school, state_code %in% c("MA", "ME", "VT")
  & visits_by_100751 >= 1))
#> [1] 108
```



## Subset function, examples

Use the `select` argument within `subset()` to keep selected variables

► syntax: `select = c(var_name1, var_name2, ..., var_name_n)`

Subset all CA public high schools that are at least 50% Latinx **AND** only keep variables `name` and `address`

```
subset(x = df_school, subset = school_type == "public" & pct_hispanic >= 50  
      & state_code == "CA", select = c(name, address))
```

```
#>                                                                 name  
#> 1254                                                                Tustin High  
#> 1301                                                                Bell Gardens High  
#> 1309                                                                Santa Ana High  
#> 1332                                                                Warren High  
#> 1336                                                                Hollywood Senior High  
#> 1337                                                                Venice Senior High  
#> 1345                                                                Sequoia High  
#> 1360                                                                Santa Barbara Senior High  
#> 1361                                                                Santa Paula High  
#> 1362                                                                Azusa High  
#> 1363                                                                Alta Loma High  
#> 1373                                                                Garden Grove High  
#> 1374                                                                Mt. Eden High  
#> 1375                                                                Ocean View High  
#> 1376                                                                Inglewood High  
#> 1378                                                                Bell Senior High  
#> 1379                                                                El Camino High  
#> 1380                                                                Hemet High
```

## Subset function, examples

Combine `subset()` with assignment (`<-`) to create a new data frame

Create a new data frame of all CA public high schools that are at least 50% Latinx  
**AND** only keep variables `name` and `address`

```
df_school_v2 <- subset(df_school, school_type == "public" & pct_hispanic >= 50  
  & state_code == "CA", select = c(name, address))
```

```
head(df_school_v2, n=5)
```

```
#>           name           address  
#> 1254      Tustin High  1171 El Camino Real  
#> 1301   Bell Gardens High  6119 Agra St.  
#> 1309   Santa Ana High   520 W. Walnut  
#> 1332    Warren High   8141 De Palma St.  
#> 1336 Hollywood Senior High 1521 N. Highland Ave.
```

```
nrow(df_school_v2)
```

```
#> [1] 713
```

# Student Exercises

Using `subset()` from base R:

1. Create a new dataframe by extracting the columns `instnm`, `event_date`, `event_type` from `df_event` data frame. And show what columns (variables) are in the newly created dataframe.
2. Create a new dataframe from the `df_school` data frame that includes out-of-state public high schools with 50%+ Latinx student enrollment that received at least one visit by the University of California Berkeley (`var=visits_by_110635`). And count the number of observations.
3. Count the number of public schools from CA, FL or MA that received one or two visits from UC Berkeley from the `df_school` data frame.
4. Subset all public out-of-state high schools visited by University of California Berkeley that enroll at least 50% Black students, and only keep variables `state_code`, `name` and `zip_code`.

## Solution to Student Exercises

### Solution to 1

```
df_event_br <- subset(df_event, select=c(instnm, event_date, event_type))  
names(df_event_br)  
#> [1] "instnm"      "event_date" "event_type"
```

### Solution to 2

```
df_school_br <- subset(df_school, state_code != "CA" & school_type == "public"  
                        & pct_hispanic >= 50 & visits_by_110635 >=1 )  
nrow(df_school_br)  
#> [1] 10
```

### Solution to 3

```
nrow(subset(df_school, state_code %in% c("CA", "FL", "MA")  
            & school_type == "public" & visits_by_110635 %in% c(1,2) ))  
#> [1] 246
```

## Solution to Student Exercises

### Solution to 4

```
subset(df_school, school_type == "public" & state_code != "CA"
      & visits_by_110635 >= 1 & pct_black >= 50,
      select = c(state_code, name, zip_code))
```

#>	state_code	name	zip_code
#> 4523	GA	Grady High School	30309
#> 8519	MD	Frederick Douglass High	20772
#> 9666	MN	DOWNTOWN CAMPUS	55403
#> 10681	MS	MURRAH HIGH SCHOOL	39202
#> 14569	OH	Shaker Hts High School	44120
#> 14632	OH	Cleveland Heights High School	44118
#> 17142	SC	Spring Valley High	29229
#> 17143	SC	Richland Northeast High	29223
#> 17623	TN	Soulsville Charter School	38106
#> 17624	TN	KIPP Memphis Collegiate High School	38108

## 4 Creating variables

## Create new data frame based on `df_school_all`

Data frame `df_school_all` has one obs per US high school and then variables identifying number of visits by particular universities

```
load(url("https://github.com/ozanj/rclass/raw/master/data/recruiting/recruit_schools.csv"),
names(df_school_all))
```

```
#> [1] "state_code"      "school_type"      "necessch"
#> [4] "name"            "address"          "city"
#> [7] "zip_code"        "pct_white"        "pct_black"
#> [10] "pct_hispanic"    "pct_asian"        "pct_amerindian"
#> [13] "pct_other"       "num_fr_lunch"     "total_students"
#> [16] "num_took_math"   "num_prof_math"    "num_took_rola"
#> [19] "num_prof_rola"   "avgmedian_inc_2564" "latitude"
#> [22] "longitude"       "visits_by_196097" "visits_by_186380"
#> [25] "visits_by_215293" "visits_by_201885" "visits_by_181464"
#> [28] "visits_by_139959" "visits_by_218663" "visits_by_100751"
#> [31] "visits_by_199193" "visits_by_110635" "visits_by_110653"
#> [34] "visits_by_126614" "visits_by_155317" "visits_by_106397"
#> [37] "visits_by_149222" "visits_by_166629" "total_visits"
#> [40] "inst_196097"     "inst_186380"      "inst_215293"
#> [43] "inst_201885"     "inst_181464"      "inst_139959"
#> [46] "inst_218663"     "inst_100751"      "inst_199193"
#> [49] "inst_110635"     "inst_110653"      "inst_126614"
#> [52] "inst_155317"     "inst_106397"      "inst_149222"
#> [55] "inst_166629"
```

## Create new data frame based on `df_school_all`

Create new version of data frame, called `school_v2`, which we'll use to introduce how to create new variables

```
library(tidyverse) # below code use tidyverse functions and pipe operator
#> -- Attaching core tidyverse packages ----- tidyverse 2.0.0
#> v dplyr      1.1.3      v readr      2.1.4
#> v forcats    1.0.0      v stringr   1.5.0
#> v ggplot2    3.4.3      v tibble    3.2.1
#> v lubridate  1.9.3      v tidyr     1.3.0
#> v purrr      1.0.2
#> -- Conflicts ----- tidyverse_conflicts()
#> x dplyr::filter() masks stats::filter()
#> x dplyr::lag()     masks stats::lag()
#> i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all co
school_v2 <- df_school_all %>%
  select(-contains("inst_")) %>% # remove vars that start with "inst_"
  rename( # rename selected variables
    visits_by_berkeley = visits_by_110635,
    visits_by_boulder = visits_by_126614,
    visits_by_bama = visits_by_100751,
    visits_by_stonybrook = visits_by_196097,
    visits_by_rutgers = visits_by_186380,
    visits_by_pitt = visits_by_215293,
    visits_by_cinci = visits_by_201885,
    visits_by_nebraska = visits_by_181464,
    visits_by_georgia = visits_by_139959,
    visits_by_scarolina = visits_by_218663,
```



## Base R approach to creating new variables

Create new variables using assignment operator `<-` and subsetting operators `[]` and `$` to create new variables and set conditions of the input variables

Pseudo syntax: `df$newvar <- ...`

- ▶ where `...` argument is expression(s)/calculation(s) used to create new variables
- ▶ expressions can include subsetting operators and/or other base R functions

**Task:** Create measure of percent of students on free-reduced lunch

### base R approach

```
school_v2_temp <- school_v2 #create copy of dataset; not necessary
school_v2_temp$pct_fr_lunch <-
  school_v2_temp$num_fr_lunch/school_v2_temp$total_students

#investigate variable you created
str(school_v2_temp$pct_fr_lunch)
#> num [1:21301] 0.723 1 0.967 0.93 1 ...
school_v2_temp$pct_fr_lunch[1:5] # print first 5 obs
#> [1] 0.7225549 1.0000000 0.9666667 0.9303483 1.0000000
```

### tidyverse approach (with pipes)

```
school_v2_temp <- school_v2 %>%
  mutate(pct_fr_lunch = num_fr_lunch/total_students)
```

## Base R approach to creating new variables

If creating new variable based on the condition/values of input variables, basically the tidyverse equivalent of `mutate()` **with** `if_else()` or `recode()`

- ▶ Pseudo syntax: `df$newvar[logical condition]<- new value`
- ▶ `logical condition`: a condition that evaluates to `TRUE` or `FALSE`

## Base R approach to creating new variables

**Task:** Create 0/1 indicator if school has median income greater than \$100k

### tidyverse approach (using pipes)

```
school_v2_temp %>% select med_inc) %>%  
  mutate(inc_gt_100k= if_else(med_inc>100000,1,0)) %>%  
  count(inc_gt_100k) # note how NA values of med_inc treated  
#> # A tibble: 3 x 2  
#>   inc_gt_100k      n  
#>   <dbl> <int>  
#> 1         0 18632  
#> 2         1  2045  
#> 3        NA   624
```

### Base R approach

```
school_v2_temp$inc_gt_100k<-NA #initialize an empty column with NAs  
                                # otherwise you'll get warning  
school_v2_temp$inc_gt_100k[school_v2_temp$med_inc>100000] <- 1  
school_v2_temp$inc_gt_100k[school_v2_temp$med_inc<=100000] <- 0  
count(school_v2_temp, inc_gt_100k)  
#> # A tibble: 3 x 2  
#>   inc_gt_100k      n  
#>   <dbl> <int>  
#> 1         0 18632  
#> 2         1  2045  
#> 3        NA   624
```

## Creating variables

**Task:** Using data frame `wwlist` and input vars `state` and `firstgen`, create a 4-category var with following categories:

- ▶ "instate\_firstgen"; "instate\_nonfirstgen"; "outstate\_firstgen";  
"outstate\_nonfirstgen"

**tidyverse approach (using pipes)**

```
load(url("https://github.com/ozanj/rclass/raw/master/data/prospect_list/wwlist_"))
wwlist_temp <- wwlist %>%
  mutate(state_gen = case_when(
    state == "WA" & firstgen == "Y" ~ "instate_firstgen",
    state == "WA" & firstgen == "N" ~ "instate_nonfirstgen",
    state != "WA" & firstgen == "Y" ~ "outstate_firstgen",
    state != "WA" & firstgen == "N" ~ "outstate_nonfirstgen")
  )
str(wwlist_temp$state_gen)
#> chr [1:268396] NA "instate_nonfirstgen" "instate_nonfirstgen" ...
wwlist_temp %>% count(state_gen)
#> # A tibble: 5 x 2
#>   state_gen      n
#>   <chr>      <int>
#> 1 instate_firstgen 32428
#> 2 instate_nonfirstgen 58646
#> 3 outstate_firstgen 32606
#> 4 outstate_nonfirstgen 134616
#> 5 <NA>      10100
```

## Base R approach to creating new variables

**Task:** Using `wwlist` and input vars `state` and `firstgen`, create a 4-category var

### base R approach

```
wwlist_temp <- wwlist
```

```
wwlist_temp$state_gen <- NA
wwlist_temp$state_gen[wwlist_temp$state == "WA"
  & wwlist_temp$firstgen == "Y"] <- "instate_firstgen"
wwlist_temp$state_gen[wwlist_temp$state == "WA"
  & wwlist_temp$firstgen == "N"] <- "instate_nonfirstgen"
wwlist_temp$state_gen[wwlist_temp$state != "WA"
  & wwlist_temp$firstgen == "Y"] <- "outstate_firstgen"
wwlist_temp$state_gen[wwlist_temp$state != "WA"
  & wwlist_temp$firstgen == "N"] <- "outstate_nonfirstgen"
```

```
str(wwlist_temp$state_gen)
#> chr [1:268396] NA "instate_nonfirstgen" "instate_nonfirstgen" ...
count(wwlist_temp, state_gen)
#> # A tibble: 5 x 2
#>   state_gen      n
#>   <chr>      <int>
#> 1 instate_firstgen 32428
#> 2 instate_nonfirstgen 58646
#> 3 outstate_firstgen 32606
#> 4 outstate_nonfirstgen 134616
#> 5 <NA>          10100
```

## 5 Appendix

## 5.1 Sorting data

## Base R `sort()` for vectors

`sort()` is a base R function that sorts vectors

Syntax: `sort(x, decreasing=FALSE, ...)`

- ▶ where x is object being sorted
- ▶ By default it sorts in ascending order (low to high)
- ▶ Need to set decreasing argument to `TRUE` to sort from high to low

```
##?sort()  
x<- c(31, 5, 8, 2, 25)  
sort(x)  
#> [1] 2 5 8 25 31  
sort(x, decreasing = TRUE)  
#> [1] 31 25 8 5 2
```



## Base R `order()` for dataframes

`order()` is a base R function that sorts vectors

- ▶ Syntax: `order(..., na.last = TRUE, decreasing = FALSE)`
- ▶ where `...` are variable(s) to sort by
- ▶ By default it sorts in ascending order (low to high)
- ▶ Need to set decreasing argument to `TRUE` to sort from high to low

Descending argument only works when we want either one (and only) variable descending or all variables descending (when sorting by multiple vars)

- ▶ use `-` when you want to indicate which variables are descending while using the default ascending sorting

```
df_event[order(df_event$event_date), ]  
df_event[order(df_event$event_date, df_event$total_12), ]  
  
#sort descending via argument  
df_event[order(df_event$event_date, decreasing = TRUE), ]  
df_event[order(df_event$event_date, df_event$total_12, decreasing = TRUE), ]  
  
#sorting by both ascending and descending variables  
df_event[order(df_event$event_date, -df_event$total_12), ]
```

## Example, sorting

- ▶ Create a new dataframe from `df_events` that sorts by ascending by `event_date`, ascending `event_state`, and descending `pop_total`.

**base R** using `order()` function:

```
df_event_br1 <- df_event[order(df_event$event_date, df_event$event_state,  
                              -df_event$pop_total), ]
```